

Intro to Instrumentation

Kubernetes Community Day, SCaLEx20



What was your *first* monitoring tool?

pingdom



MRTG
MULTI ROUTER TRAFFIC GRAPHER

Nagios



Set Up

<https://github.com/paigerduty/intro-to-instrumentation>



Docker Desktop

- [MAC Intel Chip](#)
- [Mac Apple Chip](#)
- [Windows](#)
- [Linux](#)

Docker Desktop

Install Docker Desktop – the fastest way to containerize applications.

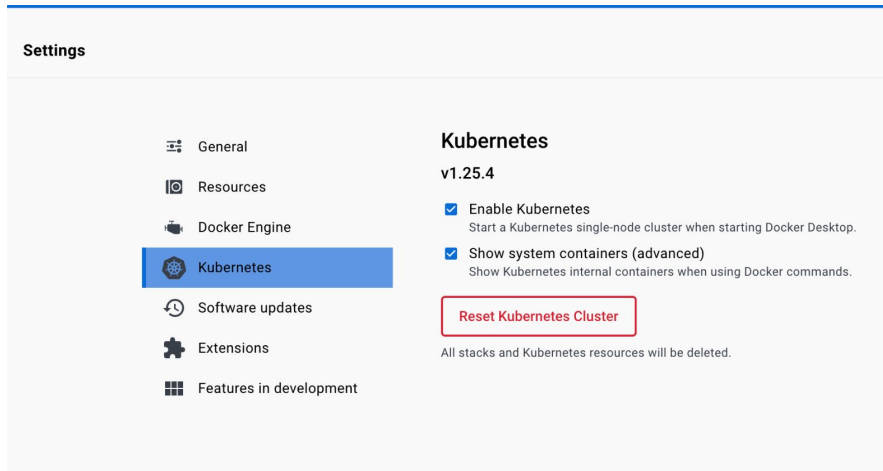
Download Docker Desktop

 Intel Chip



Docker Desktop: Enable Kubernetes

<https://docs.docker.com/desktop/kubernetes/#enable-kubernetes>



Kubectl

- [Linux](#)
- [Windows](#)
- [Mac](#)



Helm

- MacOS: brew install helm
- Windows: choco install kubernetes-helm
- Linux:

<https://helm.sh/docs/intro/install>



Observability Primer




Me + Tracing



new relic

2015-2018

SWE @ NR working
on their 1st tracing
product



2019-2021

SRE @ Lightstep (a
tracing and metrics
platform) co-founded
OTel




chronosphere

Now

OSS Instrumentation
Advocate @
Chronosphere

MFW: I think about tracing...



Never!!!



Considering it



Love it

What brought you to SCaLE?

I want to learn _____



Our agenda

- What is observability?
- What are key tracing concepts?
- What concepts do I need to use OpenTelemetry?
- How do I record data using OpenTelemetry?
- Where can I see my data?
- Next steps...



Observability Primer



What is observability?


How effectively you can understand system behavior from the outside



Why observability?

- Microservices create complex interactions.
- Failures don't exactly repeat.
- Debugging has become painstaking.
- As architectures evolve so should the approaches for making sense of interactions





Evolving visibility



View of APM /Traditional Monitoring




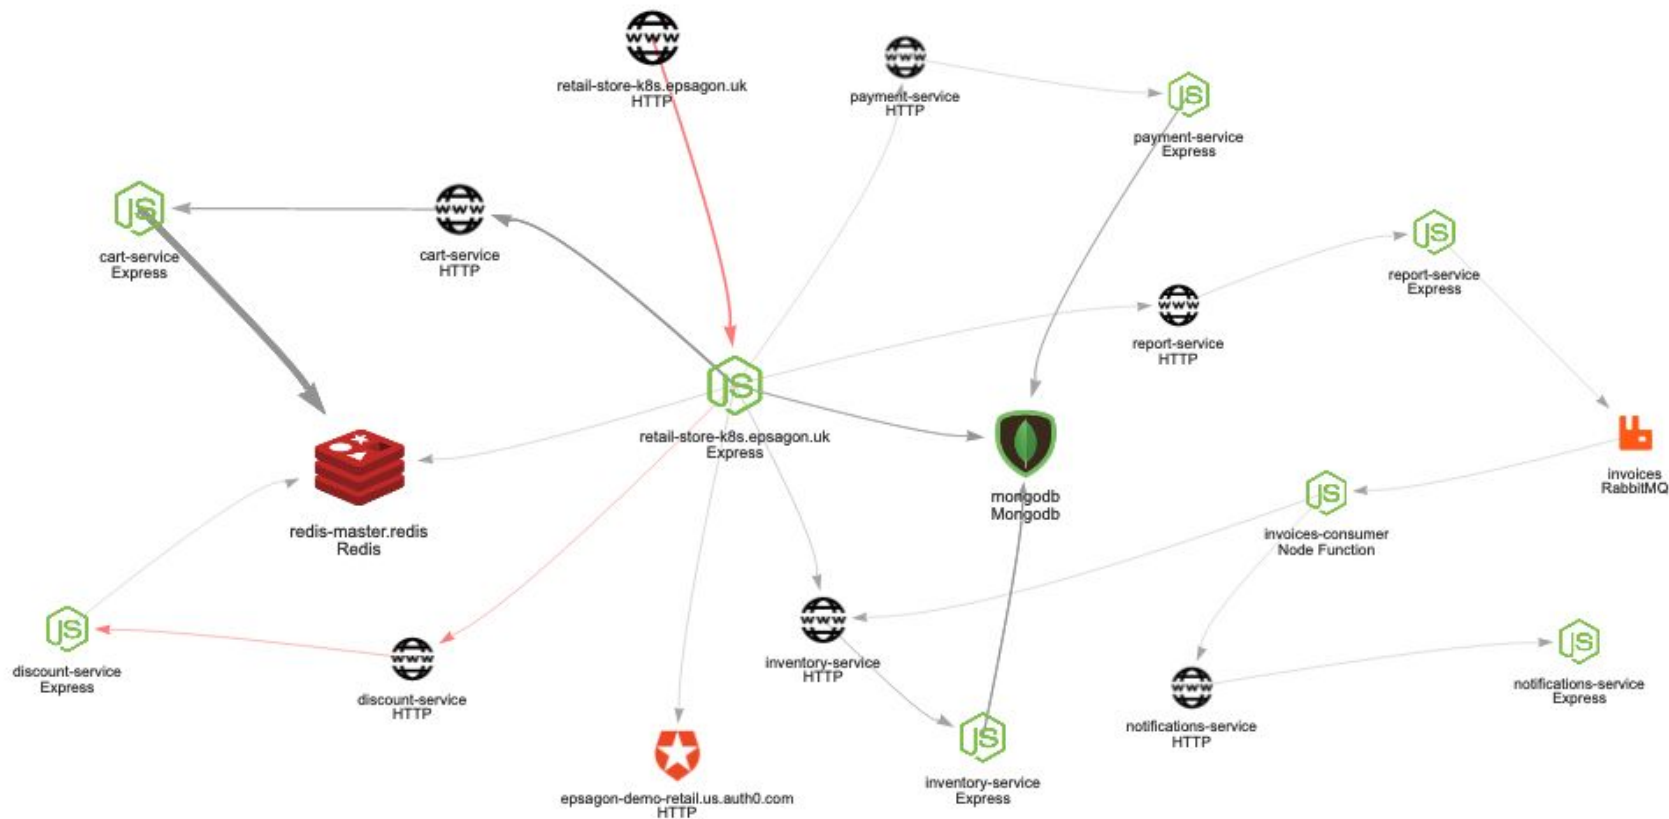
View with Distributed Tracing




Challenges

- Observe
- Troubleshoot
- Optimize







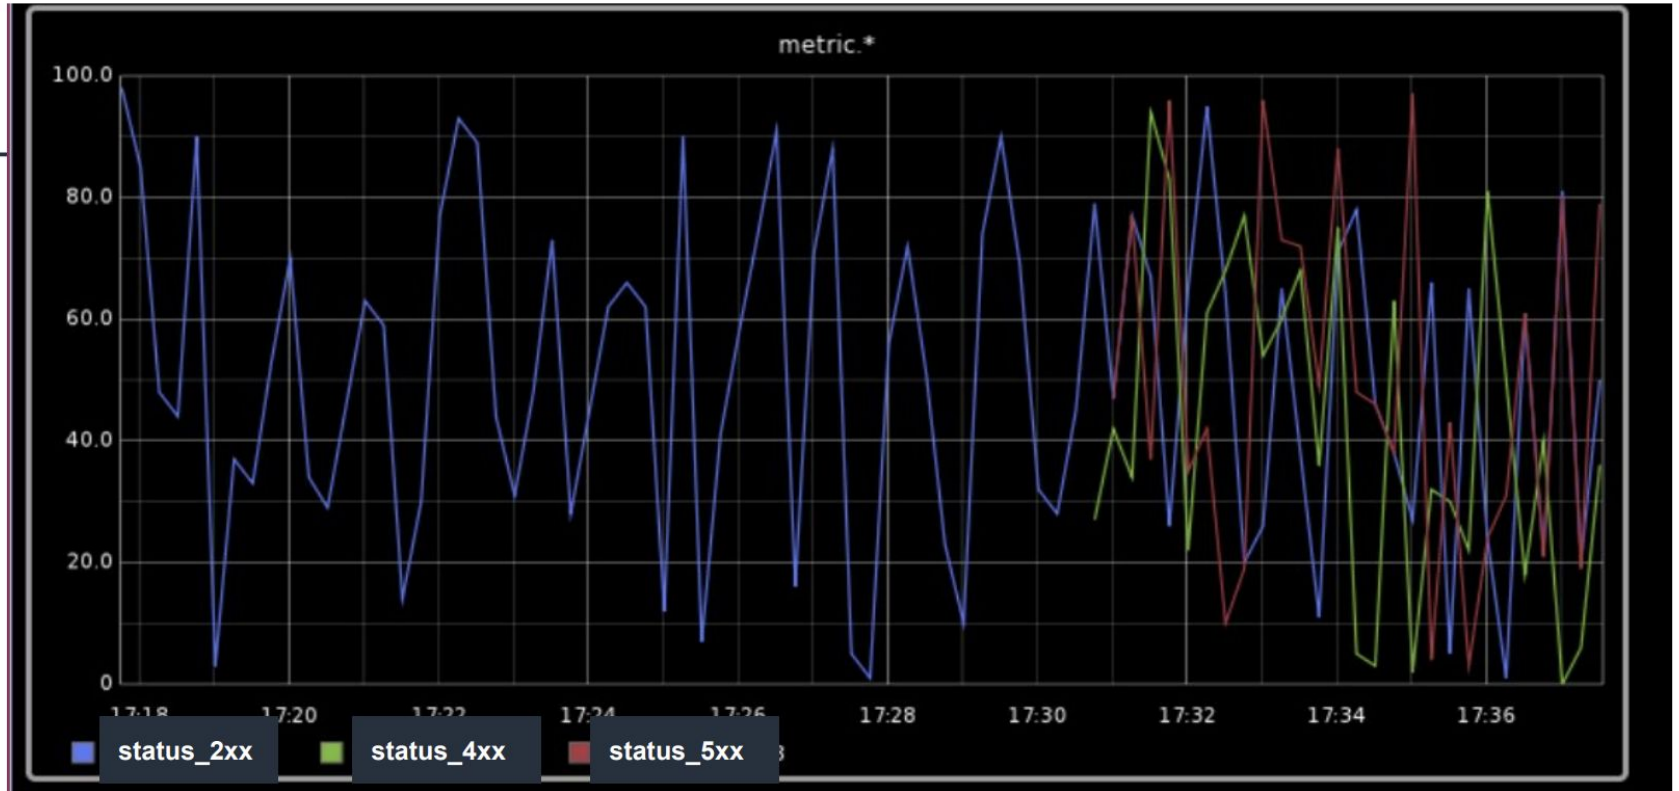
Understanding Your System

What traits did the requests that timed out at 500ms share in common? Service versions? Browser plugins? Node pool?

- Instrumentation produces data
- Querying data answers our questions.



The oncall sends you this chart showing a new spike in 400 & 500s from your service...



Telemetry aids observability

- Telemetry data **isn't** observability itself.
- Instrumenting code is *how* we get telemetry.
- Telemetry data can include traces, logs, and/or metrics.




Our agenda

- ~~What is observability?~~
- What are key tracing concepts?
- What do I know to use OpenTelemetry?
- How do I record data using OpenTelemetry?
- Where can I see my data?
- Next steps...




Tracing Concepts





trace



span



Tracing concepts in a nutshell

Trace

- Defined implicitly by its **spans**. A **trace** can be thought of as a directed acyclic graph of **spans** where the edges between **spans** are defined as parent/child relationships.



Tracing concepts in a nutshell

Span

- Represents a **single unit of work** in a system.
- Typically encapsulates: operation name, a start and finish timestamp, the parent span identifier, the span identifier, and context items.

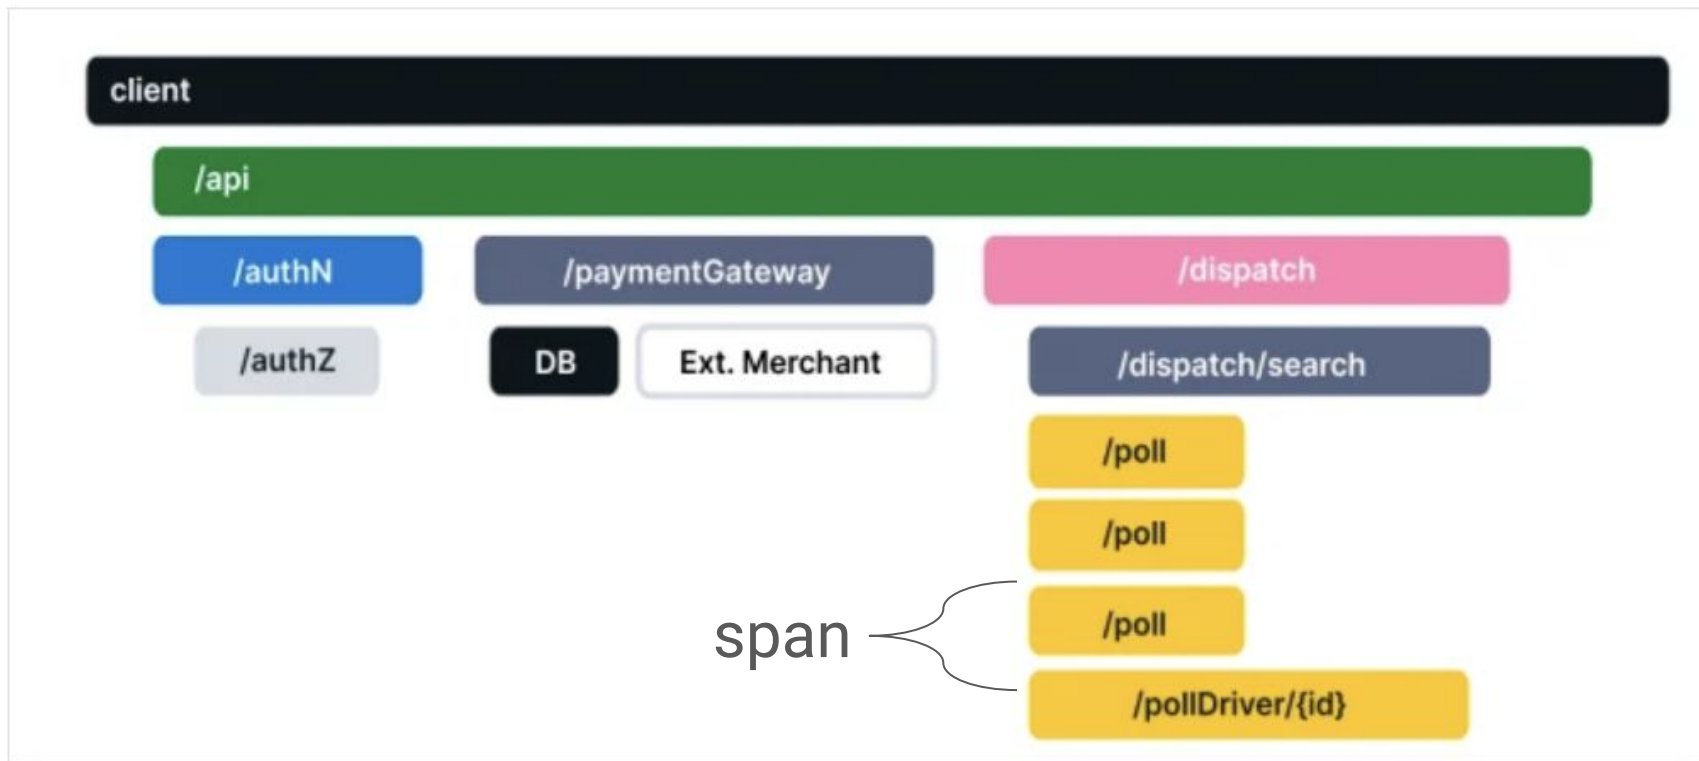


Span

```
{
  "name": "Hello-Greetings",
  "context": {
    "trace_id": "0x5b8aa5a2d2c872e8321cf37308d69df2",
    "span_id": "0x5fb397be34d26b51",
  },
  "parent_id": "0x051581bf3cb55c13",
  "start_time": "2022-04-29T18:52:58.114304Z",
  "end_time": "2022-04-29T22:52:58.114561Z",
  "attributes": {
    "http.route": "some_route1"
  },
  "events": [
    {
      "name": "hey there!",
      "timestamp": "2022-04-29T18:52:58.114561Z",
      "attributes": {
        "event_attributes": 1
      }
    },
    {
      "name": "bye now!",
      "timestamp": "2022-04-29T18:52:58.114585Z",
      "attributes": {
        "event_attributes": 1
      }
    }
  ],
}
```



trace



span



How do I implement these?

- You need an instrumentation framework!
- and a place to send the data!
- and a way to visualize the data!



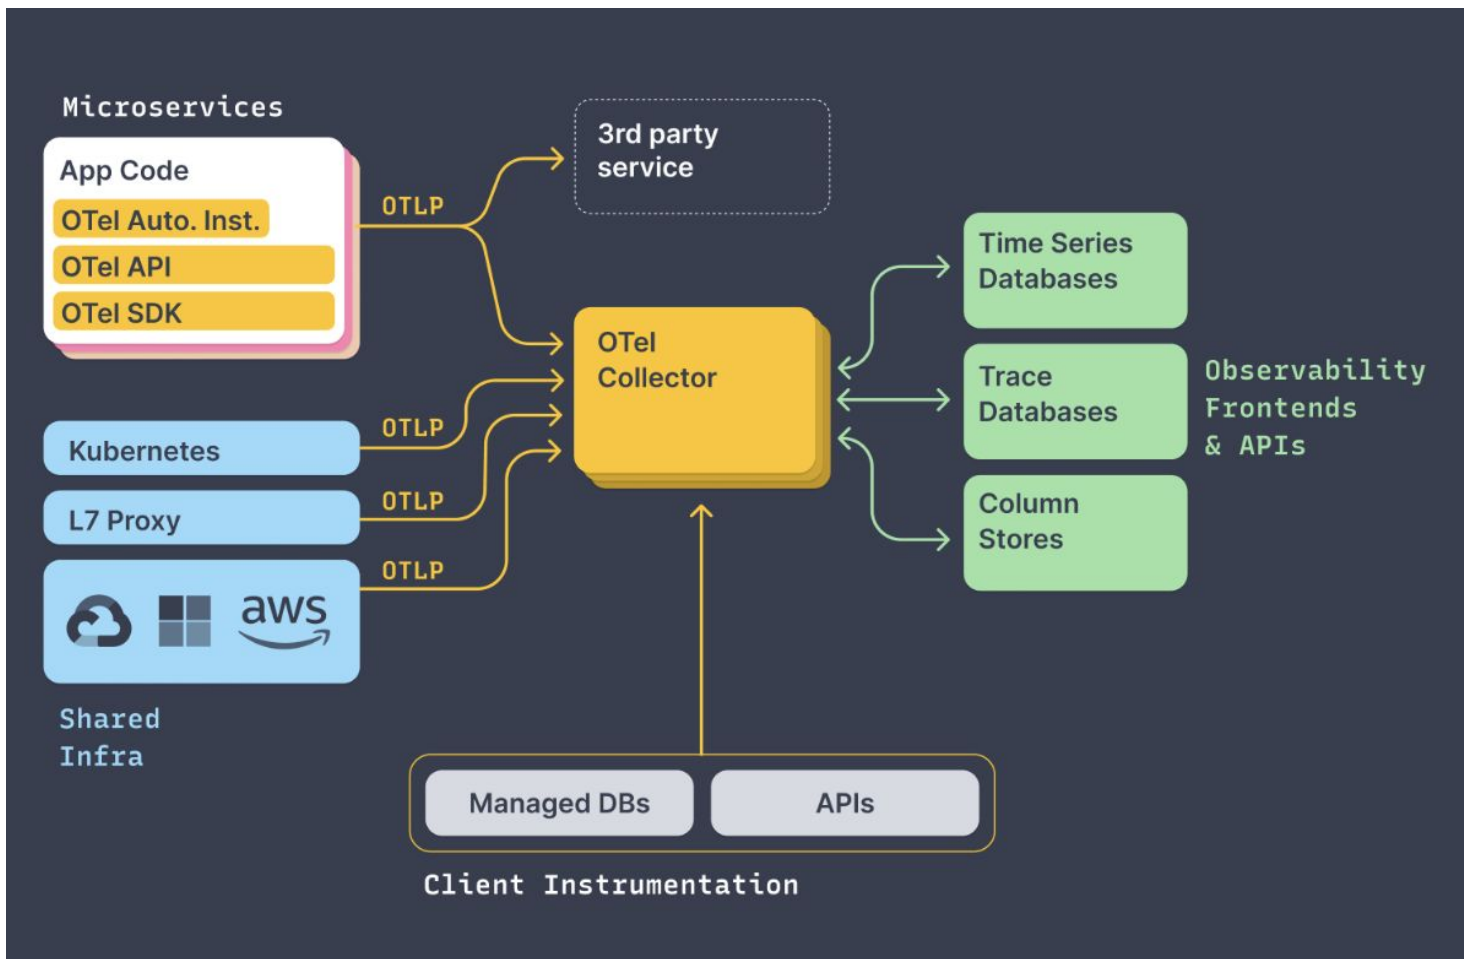
Our agenda

- ~~What is observability?~~
- ~~What are key tracing concepts?~~
- What do I know to use OpenTelemetry?
- How do I record data using OpenTelemetry?
- Where can I see my data?
- Next steps...



Hello, OpenTelemetry!









VMware
Aria™



sumo logic

UPTRACE

servicePILOT



solarwinds



Aspecto



DaoCloud



logz.io



observIQ

splunk>



Promscale



chronosphere

LogicMonitor



INSTANA



new relic



DATADOG

dynatrace



INSTANA



lumigo



elastic



APPDYNAMICS



Grafana Labs



SigNoz



Helios



Azure



honeycomb.io

Tracing

- **API:** stable, feature-freeze
- **SDK:** stable
- **Protocol:** stable
- Notes:
 - The tracing specification is now completely stable, and covered by long term support.
 - The tracing specification is still extensible, but only in a backwards compatible manner.
 - OpenTelemetry clients are versioned to v1.0 once their tracing implementation is complete.



Metrics

- **API**: stable
- **SDK**: mixed
- **Protocol**: stable
- Notes:
 - OpenTelemetry Metrics is currently under active development.
 - The data model is stable and released as part of the OTLP protocol.
 - Experimental support for metric pipelines is available in the Collector.
 - Collector support for Prometheus is under development, in collaboration with the Prometheus community.



Logging

- **API:** draft
- **SDK:** draft
- **Protocol:** stable
- Notes:
 - OpenTelemetry Logging is currently under active development.
 - The **logs data model** is released as part of the OpenTelemetry Protocol.
 - Log processing for many data formats has been added to the Collector, thanks to the donation of Stanza to the OpenTelemetry project.
 - Log appenders are currently under development in many languages. Log appenders allow OpenTelemetry tracing data, such as trace and span IDs, to be appended to existing logging systems.
 - An OpenTelemetry logging SDK is currently under development. This allows OpenTelemetry clients to ingest logging data from existing logging systems, outputting logs as part of OTLP along with tracing and metrics.
 - An OpenTelemetry logging API is not currently under development. We are focusing first on integration with existing logging systems. When metrics is complete, focus will shift to development of an OpenTelemetry logging API.



Is Your Framework Instrumented?

Think of the tech stack you use for projects or at work ...

Head to the Registry

<https://opentelemetry.io/ecosystem/registry/>

Share what you find!

Any languages or frameworks missing? Any that surprised you?



Our agenda

- ~~What is observability?~~
- ~~What are key tracing concepts?~~
- ~~What do I know to use OpenTelemetry?~~
- How do I record data using OpenTelemetry?
- Where can I see my data?
- Next steps...



Instrumenting (aka recording data)



SDKs, Exporters, and Collector Services, Oh My!

- OpenTelemetry's **SDK** implements trace & span creation.
- An **exporter** can be instantiated to send the data collected by OpenTelemetry to the backend of your choice.
- OpenTelemetry **collector** proxies data between instrumented code and backend service(s). The exporters can be reconfigured without changing instrumented code.



Manual vs Automatic Instrumentation



Automatic Instrumentation

```
@app.route("/server_request")
def server_request():
    print(request.args.get("param"))
    return "served"
```



Manual Instrumentation

```
@app.route("/server_request")
def server_request():
    with tracer.start_as_current_span(
        "server_request",
        context=extract(request.headers),
        kind=trace.SpanKind.SERVER,
        attributes=collect_request_attributes(request.environ),
    ):
        print(request.args.get("param"))
        return "served"
```



Tracer methods, & when to call

- `tracer.start_span(name, parent=, ...)`
 - This method returns a child of the specified span.
- `with tracer.start_as_current_span(name)`
 - Starts a new span, sets it to be active. Optionally, can get a reference to the span.
- `tracer.get_current_span()`
 - Used to access & add information to the current span



Span methods, & when to call

- `span.add_event(name, attributes)`
 - Adds structured annotations (e.g. "logs") about what is currently happening.
- `span.set_attribute(key, value)`
 - Adds an attribute to the current span. This may include a user id, a build id, a user-agent, etc.
- `span.end()`
 - Manually closes a span.



Add more context to traces with Span Events

- Span Events are context-aware logging.
- An `event` contains timestamped information added to a span. You can think of this as a structured log, or a way to annotate your spans with specific details about what happened along the way.
 - Contains:
 - the name of the event
 - one or more attributes
 - a timestamp



Code examples: Current Span & Span

- Get the current span
 - `span = tracer.get_current_span()`
- Update the span status
 - `span.set_status(Status(StatusCanonicalCode.UNKNOWN, error))`
- Add events
 - `span.add_event("foo", {"customer": "bar"})`
- Add attributes
 - `span.set_attribute("error", True)`



Auto-instrument



```
pip3 install opentelemetry-distro
```



```
pip3 install opentelemetry-distro
```

```
opentelemetry-bootstrap -a install
```

```
opentelemetry-instrument \  
  --traces_exporter console \  
  --metrics_exporter console \  
flask run
```




```
opentelemetry-instrument \  
  --traces_exporter console \  
  --metrics_exporter console \  
  flask run
```



```

{
  "name": "/visit",
  "context": {
    "trace_id": "0xa2f924bb3a05138820631f9913c84487",
    "span_id": "0xa79dd2bc621c37fc",
    "trace_state": "[]"
  },
  "kind": "SpanKind.SERVER",
  "parent_id": null,
  "start_time": "2023-03-09T21:19:12.509247Z",
  "end_time": "2023-03-09T21:19:12.510387Z",
  "status": {
    "status_code": "UNSET"
  },
  "attributes": {
    "http.method": "GET",
    "http.server_name": "127.0.0.1",
    "http.scheme": "http",
    "net.host.port": 5000,
    "http.host": "localhost:5000",
    "http.target": "/visit",
    "net.peer.ip": "127.0.0.1",
    "http.user_agent": "Mozilla/5.0 (Macintosh; Intel Mac OS X 10_15_7) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/111.0.0.0 Safari/537.36",
    "net.peer.port": 56044,
    "http.flavor": "1.1",
    "http.route": "/visit",
    "http.status_code": 200
  },
  "events": [],
  "links": [],
  "resource": {
    "attributes": {
      "telemetry.sdk.language": "python",
      "telemetry.sdk.name": "opentelemetry",
      "telemetry.sdk.version": "1.16.0",
      "telemetry.auto.version": "0.37b0",
      "service.name": "unknown_service"
    },
    "schema_url": ""
  }
}

```



Manually instrument



app.py

```
1  from flask import Flask
2  from opentelemetry import trace
3  from opentelemetry.sdk.trace import TracerProvider
4  from opentelemetry.sdk.trace.export import BatchSpanProcessor
5  from opentelemetry.exporter.otlp.proto.http.trace_exporter import OTLPSpanExporter
6
7
8  import random
9
10 trace.set_tracer_provider(TracerProvider())
11 trace.get_tracer_provider().add_span_processor(
12     BatchSpanProcessor(OTLPSpanExporter())
13 )
14
15 # Acquire a tracer
16 tracer = trace.get_tracer(__name__)
17
```

app.py

```
31
32 @app.route("/rolldice")
33 def roll_dice():
34     return str(do_roll())
35
36 def do_roll():
37     with tracer.start_as_current_span("do_roll") as rollspan:
38         result = random.randint(1, 6)
39         rollspan.set_attribute("roll.value", result)
40     return result
41
42 if __name__ == "__main__":
43     app.run(host="0.0.0.0", port=8000, debug=True)
44
```



```
pip3 install -r requirements.txt
```




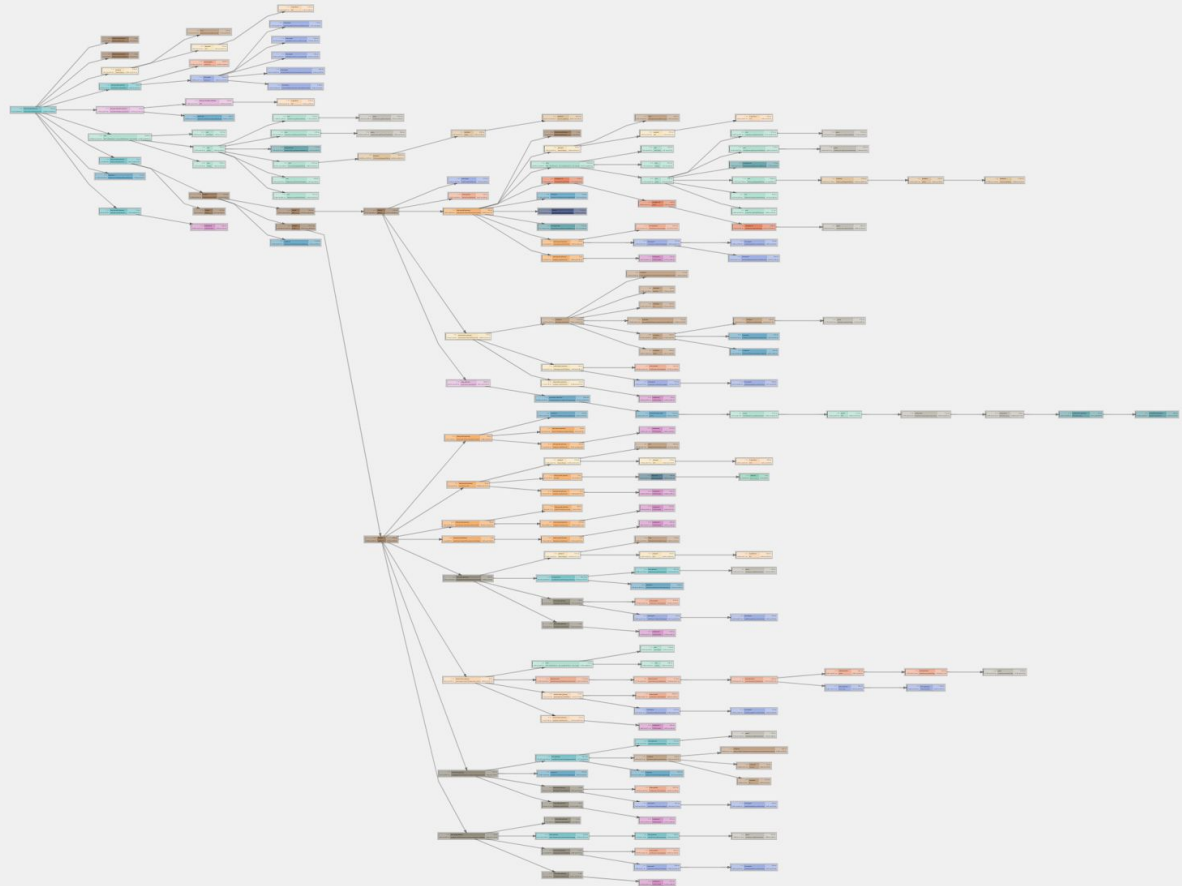
OTEL_SERVICE_NAME=appy flask run



Jaeger!



Trace Start **September 4 2019, 18:47:33.503** | Duration **618ms** | Services **32** | Depth **16** | Total Spans **352**



```
docker run -d --name jaeger \  
-e COLLECTOR_OTLP_ENABLED=true \  
-p 16686:16686 \  
-p 4317:4317 \  
-p 4318:4318 \  
jaegertracing/all-in-one:1.42
```



Go look for your trace!

The Jaeger visualization URL is at (notice the port number):

`http://localhost:16686/search`

Put in your `SERVICE_NAME` value into the service name, and search for your recent traces!



Our agenda

- ~~What is observability?~~
- ~~What are key tracing concepts?~~
- ~~What do I know to use OpenTelemetry?~~
- ~~How do I record data using OpenTelemetry?~~
- ~~Where can I see my data?~~
- Next steps...



Community

- [Documentation](#)
- [GitHub](#)
- [CNCF Slack](#)
- [SIG meetings](#)



TIL _____

